

## 1 Assignment 1: Introduction to TensorFlow and Keras

### Objective:

Install TensorFlow and Keras

Verify the installation

Load a dataset

Build and train a simple model

```
[1]: # Step 1: Install TensorFlow (uncomment if needed)
# !pip install tensorflow
```

```
[2]: # Step 2: Import Libraries
import tensorflow as tf
from tensorflow import keras
import numpy as np
import matplotlib.pyplot as plt
```

```
[3]: # Step 3: Print TensorFlow and Keras Versions
print("TensorFlow version:", tf.__version__)
print("Keras version:", keras.__version__)
```

TensorFlow version: 2.18.0

Keras version: 3.8.0

```
[4]: # Step 4: Load MNIST dataset
mnist = keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
print("Train shape:", x_train.shape)
print("Test shape:", x_test.shape)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 0s

0us/step

Train shape: (60000, 28, 28)

Test shape: (10000, 28, 28)

```
[5]: # Step 5: Normalize Data
x_train = x_train / 255.0
x_test = x_test / 255.0
```

```
[6]: #Step 6: Build Model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28,28)),
    keras.layers.Dense(64, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead. super().\_init\_(\*kwargs)

```
[7]: # Step 7: Compile Model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
[8]: # Step 8: Train Model
model.fit(x_train, y_train, epochs=2)
```

Epoch 1/2

1875/1875 7s 3ms/step -

accuracy: 0.8612 - loss: 0.4987

Epoch 2/2

1875/1875 5s 3ms/step -

accuracy: 0.9544 - loss: 0.1571

[8]: <keras.src.callbacks.history.History at 0x7e434222ea90>

```
[9]: # Step 9: Evaluate Model
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Test Accuracy:", test_acc)
```

313/313 1s 4ms/step -

accuracy: 0.9594 - loss: 0.1346 Test

Accuracy: 0.9635000228881836

This notebook introduces TensorFlow and Keras by building a simple neural network to classify handwritten digits from the MNIST dataset. It includes steps for loading and normalizing data, defining and training a model, and evaluating its accuracy. The model achieves predictions using a basic feedforward architecture. Visualizations help verify both data and model performance.

**1.1 Instruction to Student Instructions:**

- Run each cell step by step.
- Write 2-3 lines **explaining what happens in each step.**
- Take a screenshot of the final accuracy output.
- Download your notebook as .ipynb and as PDF.
- Submit both files.